

Example workloads for networking applications

Pravin Shinde

Contents

| | | |
|----------|--|----------|
| 1 | Workloads | 5 |
| 1.1 | The Big picture | 5 |
| 1.1.1 | Application requirements: | 5 |
| 1.1.2 | Hardware features: | 6 |
| 1.1.3 | Questions for given application: | 6 |
| 1.1.4 | Application: DNS Server (eg: Bind) | 6 |
| 1.1.5 | Application: HTTP server (eg: apache) | 7 |
| 1.1.6 | Application: Database server (eg: mysql!!) | 8 |
| 1.1.7 | Application: NFS filesystem client | 8 |
| 1.1.8 | Application: MPI application | 9 |
| 1.1.9 | Application: Web crawler | 9 |
| 1.2 | test diagrams | 10 |

Chapter 1

Workloads

1.1 The Big picture

Applications have some requirements, and hardware has some features

1.1.1 Application requirements:

Following are the possible application requirements.

1. Number of listen ports
2. Number of incoming clients
3. Avg number of packets in each connection
4. Connection lifetime
5. Incoming traffic expected
6. Outgoing traffic expected
7. Outgoing active connections expected
8. Protocol type
9. Number of cores involved
10. Memory locations
11. Low latency preference
12. High throughput preference
13. Low jitter requirements

1.1.2 Hardware features:

1. Dedicated RX queues
2. Hardware filters
3. Large receive offload
4. Large send offload
5. TCP offload

1.1.3 Questions for given application:

1. Does it make sense to use these hardware features?
2. Which hardware features should be used?
3. How they should be configured?
4. How they should be used?

How do I use given resources to best suit application requirements?

1.1.4 Application: DNS Server (eg: Bind)

Requirements:

1. UDP protocol
2. Single listen port
3. Large number of incoming clients
4. Single packet request/response
5. Preference to throughput
6. Load balancing with multiple cores

Ideal hardware setup

1. One dedicated RX queue for each load balancing core.
2. Hardware filter for (protocol, destination IP, destination port) to separate packets for this particular application.
3. Ideal : Give each separated packet to next core in round-robin fashion

Alternate setup: 1

1. Hash (source ip, source port) (assuming uniform distribution of hash values)
2. Use hash to select one of the dedicated RX queues.

Alternate setup: 2

1. Use one RX queue and one filter (protocol, destination IP, destination port) for this application, and let all the cores share the same RX queue. (not a good idea due to contention on updating RX)
-

1.1.5 Application: HTTP server (eg: apache)

A server responding with small sized static pages

Requirements:

1. TCP protocol
2. Single listen port
3. Preference to throughput
4. Large number of incoming clients
5. Small request, larger response.
6. Load balancing with multiple cores

Ideal hardware setup

1. One dedicated RX queue for each load balancing core.
2. Hardware filter for (protocol, destination IP, destination port)
3. Ideal : Give each new connection to next core in round-robin fashion

Alternate setup: 1

1. Hash (source ip, source port)
2. Use hash to select one of the dedicated RX queues.

Alternate setup: 2

1. Use syn filter to separate syn packets
2. Give all syn packets to load balancer
3. Load balancer will distribute them in round-robin manner
4. Insert new flow directing filter to ensure rest of the packet of this connection goes directly to proper core.

1.1.6 Application: Database server (eg: mysql!!)

A server handling small number of clients with large number of queries

Requirements:

1. TCP protocol
2. Single listen port
3. Preference to throughput
4. small number of incoming clients
5. Small request, large response.
6. Load balancing with ??

Not complete and useful yet!,
get more accurate data about
this setup

1.1.7 Application: NFS filesystem client

A kernel code which connects to NFS server and gets the contents of files based on application requests.

Requirements:

1. UDP protocol
2. Single connect port (outgoing connection)
3. Preference to throughput
4. Small request, large response. (reading data)
5. Load balancing: Increase number of kernel threads doing IO over NFS. The queries and responses are marked by RPC transaction-IDs which can be used to map the responses to proper kernel thread.

Ideal hardware setup

1. One dedicated RX queue for each load balancing core.
2. Hardware filter for (protocol, destination IP, destination port)
3. Give each packet to proper kernel thread based on RPC transaction ID.

Alternate setup: 1

If there is only one application

1. Use hash(source IP, source port) to select the destination core.

1.1.8 Application: MPI application

A class of scientific applications which communicate with each other using runtimes like MPI and perform some computation in distributed fashion.

Requirements:

1. TCP Protocol
2. Small messages
3. Large number of messages
4. Preference to low latency
5. Scalability with number of nodes
6. Long connection lifetime. (Assuming all messages are using same channel established once per node)

Ideal hardware setup

1. How many filters?

1.1.9 Application: Web crawler**Requirements:**

1. HTTP/TCP Protocol
2. Large number of outgoing connections
3. Large incoming data
4. Relatively small connection lifetime (HTTP requests)
5. Scale by adding more cores running same application with different targets
6. Preference to throughput

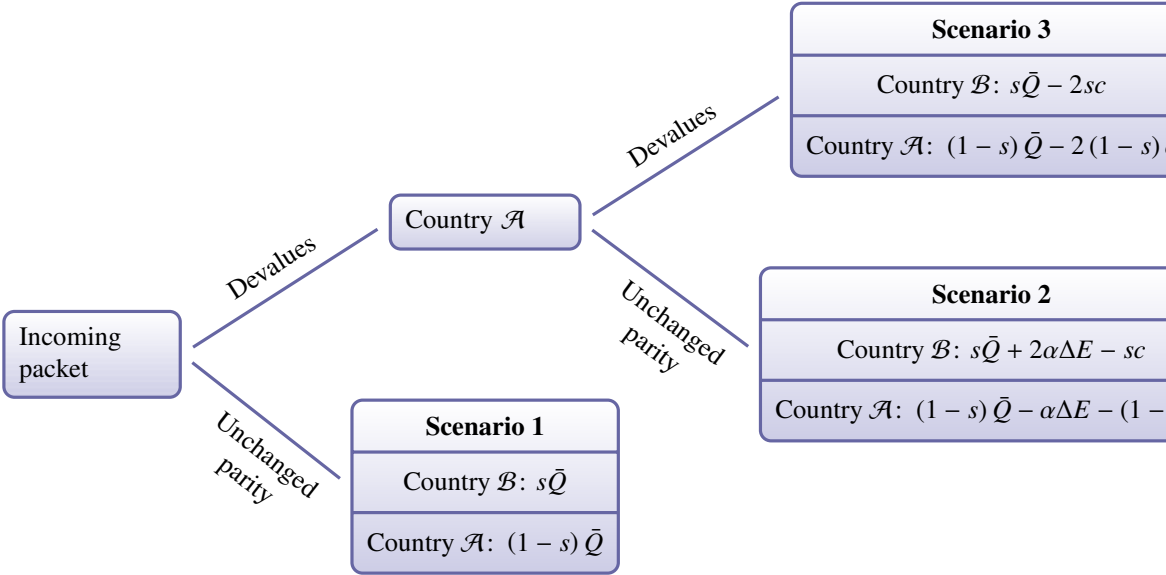
Complete the requirements of this application

Ideal hardware setup

- 1. How many filters?

1.2 test diagrams

test section



test citation [1]

Bibliography

- [1] BAUMANN, A., BARHAM, P., DAGAND, P.-E., HARRIS, T., ISAACS, R., PETER, S., ROSCOE, T., SCHÜPBACH, A., AND SINGHANIA, A. The multikernel: a new os architecture for scalable multicore systems. In *Proceedings of the 22nd Symposium on Operating Systems Principles* (New York, NY, USA, 2009), SOSP '09, ACM, pp. 29–44.