Example workloads for networking applications

Pravin Shinde

Contents

Chapter 1

Workloads

1.1 The Big picture

Applications have some requirements, and hardware has some features

1.1.1 Application requirements:

Following are the possible application requirements.

1. Number of listen ports

2. Number of incoming clients

3. Avg number of packets in each connection

4. Connection lifetime

5. Incoming traffic expected

6. Outgoing traffic expected

7. Outgoing active connections expected

8. Protocol type

9. Number of cores involved

10. Memory locations

11. Low latency preference

12. High throughput preference

13. Low jitter requirements

1.1.2   Hardware features:

1. Dedicated RX queues

2. Hardware filters

3. Large receive offload

4. Large send offload

5. TCP offload

1.1.3   Questions for given application:

1. Does it make sense to use these hardware features?

2. Which hardware features should be used?

3. How they should be configured?

4. How they should be used?

   How do I use given resources to best suit application requirements?
   ———————————————

1.1.4   Application: DNS Server (eg: Bind)

Requirements:

1. UDP protocol

2. Single listen port

3. Large number of incoming clients

4. Single packet request/response

5. Preference to throughput

6. Load balancing with multiple cores

Ideal hardware setup

1. One dedicated RX queue for each load balancing core.

2. Hardware filter for (protocol, destination IP, destination port) to separate packets
   for this particular application.

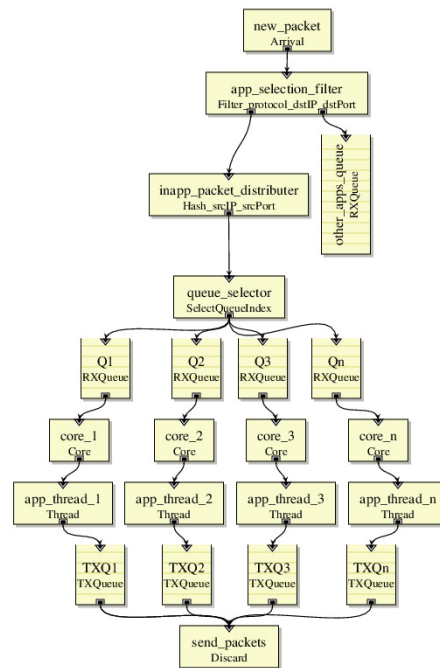3. Ideal : Give each separated packet to next core in round-robin fashion

Figure 1.1: Ideal configuration for DNS server

Alternate setup: 1

1. Hash (source ip, source port) (assuming uniform distribution of hash values)

2. Use hash to select one of the dedicated RX queues.

Alternate setup: 2

1. Use one RX queue and one filter (protocol, destination IP, destination port) for this application, and let all the cores share the same RX queue. (not a good idea due to contention on updating RX)

———————————————

1.1.5   Application: HTTP server (eg: apache)

A server responding with small sized static pages
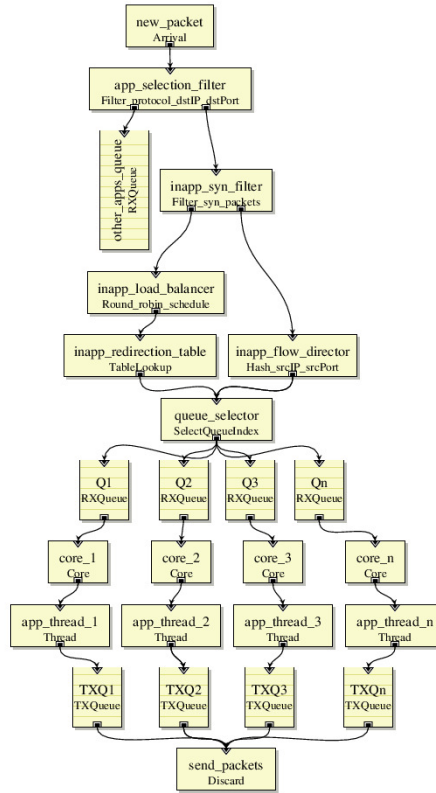
Requirements:

1. TCP protocol

Figure 1.2: Ideal configuration for HTTP server

2. Single listen port

3. Preference to throughput

4. Large number of incoming clients

5. Small request, larger response.

6. Load balancing with multiple cores

Ideal hardware setup

1. One dedicated RX queue for each load balancing core.

2. Hardware filter for (protocol, destination IP, destination port)

3. Ideal : Give each new connection to next core in round-robin fashion

Alternate setup: 1

1. Hash (source ip, source port)

2. Use hash to select one of the dedicated RX queues.

Alternate setup: 2

1. Use syn filter to separate syn packets

2. Give all syn packets to load balancer

3. Load balancer will distribute them in round-robin manner

4. Insert new flow directing filter to ensure rest of the packet of this connection goes directly to proper core.

1.1.6   Application: Web crawler

Requirements:

1. HTTP/TCP Protocol

2. Large number of outgoing connections

3. Large incoming data

4. Relatively small connection lifetime (HTTP requests)

5. Scale by adding more cores running same application with different targets

6. Preference to throughput

Ideal hardware setup

1. A dedicated RX queue for every core involved in load-balancing

2. A dedicated hardware filter for every outgoing connection

3. Filtering based on ($protocol, src-ip, src-port, dest-IP, dest-port$)

4. Dedicated queue for each load-balancing core

1.1.7   Application: Web Proxy

Requirements:

1. HTTP/TCP Protocol

2. Single listen port

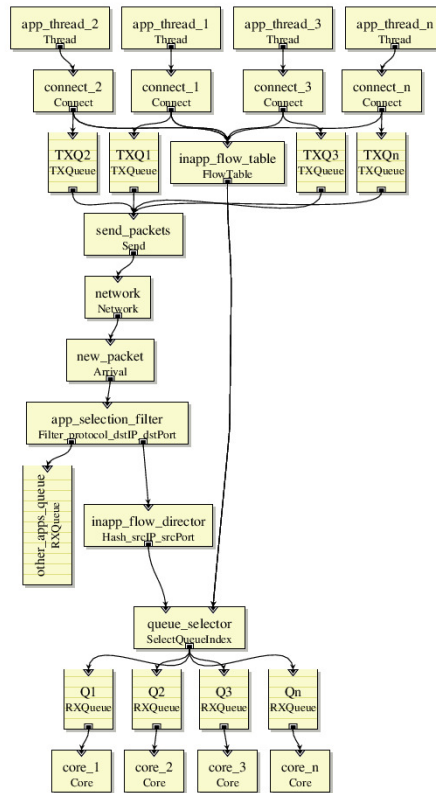3. Large number of incoming connections

Figure 1.3: Ideal configuration for Web crawler

4. Large number of outgoing connections

5. Relatively small connection lifetime for outgoing connections.

6. Larger connection lifetime for incoming connections (client connections)

7. Scaling??

8. Preference to latency and throughput (not sure)

9. With high probability, Incoming and outgoing connections are on different interfaces

10. Examples: Squid

Ideal hardware setup

1. One dedicated hardware filter for incoming connections

2. Filtering based on (protocol, destination IP, destination port)

3. Hashing to load balance connections hash(source IP, source port)

4. Dedicated queue for each load-balancing core

5. Distribution of connections across cores in round-robin fashion.

6. Does not make much sense to have a dedicated queues for outgoing connections.

1.1.8   Application: NFS filesystem client

A kernel code which connects to NFS server and gets the contents of files based on application requests.

Requirements:

1. UDP protocol

2. Single connect port (outgoing connection)

3. Preference to throughput

4. Small request, large response. (reading data)

5. Load balancing: Increase number of kernel threads doing IO over NFS. The queries and responses are marked by RPC transaction-IDs which can be used to map the responses to proper kernel thread.

Ideal hardware setup

1. One dedicated RX queue for each load balancing core.

2. Hardware filter for (protocol, destination IP, destination port)

3. Give each packet to proper kernel thread based on RPC transaction ID.

Alternate setup: 1

If there is only one application

1. Use hash(source IP, source port) to select the destination core.

1.1.9   Application: MPI application

A class of scientific applications which communicate with each other using runtimes like MPI and perform some computation in distributed fashion.

Requirements:

1. TCP Protocol

2. Small messages

3. Large number of messages

4. Preference to low latency

5. Scalability with number of nodes

6. Long connection lifetime. (Assuming all messages are using same channel established once per node)

Ideal hardware setup

1. How many filters?

1.1.10   Application: Database server (eg: mysql!!)

A server handling small number of clients with large number of queries

Requirements:

1. TCP protocol

2. Single listen port

3. Preference to throughput

4. small number of incoming clients

5. Small request, large response.

6. Load balancing with ??

1.1.11   Application: Firewall

Requirements:

1. Protocol??

Ideal hardware setup

1. How many filters?

1.1.12   Application: Intrusion Detection System

Requirements:

1. Protocol??

Ideal hardware setup

1. How many filters?

1.1.13   Application:

Requirements:

1. Protocol??

Ideal hardware setup

1. How many filters?

1.1.14   Application:

Requirements:

1. Protocol??

Ideal hardware setup

1. How many filters?

test citation [1]

Bibliography

[1] BAUMANN, A., BARHAM, P., DAGAND, P.-E., HARRIS, T., ISAACS, R., PETER, S., ROSCOE, T., SCHÜPBACH, A., AND SINGHANIA, A. The multikernel: a new os architecture for scalable multicore systems. In *Proceedings of the 22nd Symposium on Operating Systems Principles* (New York, NY, USA, 2009), SOSP '09, ACM, pp. 29–44.